

Loan Prediction

November 18, 2018

1 Loan Prediction

1.1 Problem

- A Company wants to automate the loan eligibility process (real time) based on customer detail provided while filling online application form. These details are Gender, Marital Status, Education, Number of Dependents, Income, Loan Amount, Credit History and others. To automate this process, they have given a problem to identify the customers segments, those are eligible for loan amount so that they can specifically target these customers. Here they have provided a data set.

1.2 Data

- Variable Descriptions:

Variable	Description
Loan_ID	Unique Loan ID
Gender	Male/ Female
Married	Applicant married (Y/N)
Dependents	Number of dependents
Education	Applicant Education (Graduate/ Under Graduate)
Self_Employed	Self employed (Y/N)
ApplicantIncome	Applicant income
CoapplicantIncome	Coapplicant income
LoanAmount	Loan amount in thousands
Loan_Amount_Term	Term of loan in months
Credit_History	credit history meets guidelines
Property_Area	Urban/ Semi Urban/ Rural
Loan_Status	Loan approved (Y/N)

- Rows: 615
- Source: Datahack
- Jupyter Notebook: [Github - Parth Shandilya](#)

```
In [45]: # Importing Library
import pandas as pd
```

```

import numpy as np
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder

# Reading the training dataset in a dataframe using Pandas
df = pd.read_csv("train.csv")

# Reading the test dataset in a dataframe using Pandas
test = pd.read_csv("test.csv")

```

In [48]: # First 10 Rows of training Dataset

```
df.head(10)
```

```

Out[48]:
   Loan_ID  Gender  Married  Dependents  Education  Self_Employed  \
0  LP001002   Male      No           0    Graduate             No
1  LP001003   Male     Yes           1    Graduate             No
2  LP001005   Male     Yes           0    Graduate             Yes
3  LP001006   Male     Yes           0  Not Graduate             No
4  LP001008   Male     No           0    Graduate             No
5  LP001011   Male     Yes           2    Graduate             Yes
6  LP001013   Male     Yes           0  Not Graduate             No
7  LP001014   Male     Yes           3+    Graduate             No
8  LP001018   Male     Yes           2    Graduate             No
9  LP001020   Male     Yes           1    Graduate             No

   ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term  \
0             5849              0.0         NaN         360.0
1             4583            1508.0        128.0         360.0
2             3000              0.0         66.0         360.0
3             2583            2358.0        120.0         360.0
4             6000              0.0        141.0         360.0
5             5417            4196.0        267.0         360.0
6             2333            1516.0         95.0         360.0
7             3036            2504.0        158.0         360.0
8             4006            1526.0        168.0         360.0
9            12841           10968.0        349.0         360.0

   Credit_History  Property_Area  Loan_Status
0              1.0           Urban           Y
1              1.0           Rural           N
2              1.0           Urban           Y
3              1.0           Urban           Y
4              1.0           Urban           Y
5              1.0           Urban           Y
6              1.0           Urban           Y
7              0.0       Semiurban           N
8              1.0           Urban           Y
9              1.0       Semiurban           N

```

```
In [206]: # Store total number of observation in training dataset
df_length = len(df)

# Store total number of columns in testing data set
test_col = len(test.columns)
```

2 Understanding the various features (columns) of the dataset.

```
In [50]: # Summary of numerical variables for training data set
```

```
df.describe()
```

```
Out[50]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term \
count	614.000000	614.000000	592.000000	600.00000
mean	5403.459283	1621.245798	146.412162	342.00000
std	6109.041673	2926.248369	85.587325	65.12041
min	150.000000	0.000000	9.000000	12.00000
25%	2877.500000	0.000000	100.000000	360.00000
50%	3812.500000	1188.500000	128.000000	360.00000
75%	5795.000000	2297.250000	168.000000	360.00000
max	81000.000000	41667.000000	700.000000	480.00000

	Credit_History
count	564.000000
mean	0.842199
std	0.364878
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000

1. For the non-numerical values (e.g. Property_Area, Credit_History etc.), we can look at frequency distribution to understand whether they make sense or not.

```
In [51]: # Get the unique values and their frequency of variable Property_Area
```

```
df['Property_Area'].value_counts()
```

```
Out[51]: Semiurban    233
Urban              202
Rural              179
Name: Property_Area, dtype: int64
```

2. Understanding Distribution of Numerical Variables

- ApplicantIncome
- LoanAmount

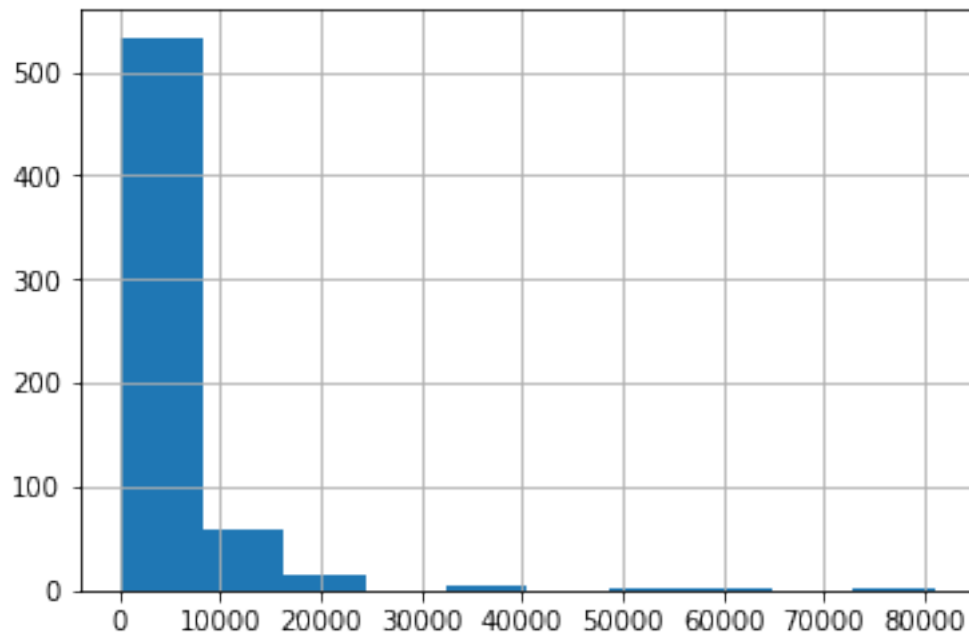
```
In [53]: # Box Plot for understanding the distributions and to observe the outliers.
```

```
%matplotlib inline
```

```
# Histogram of variable ApplicantIncome
```

```
df['ApplicantIncome'].hist()
```

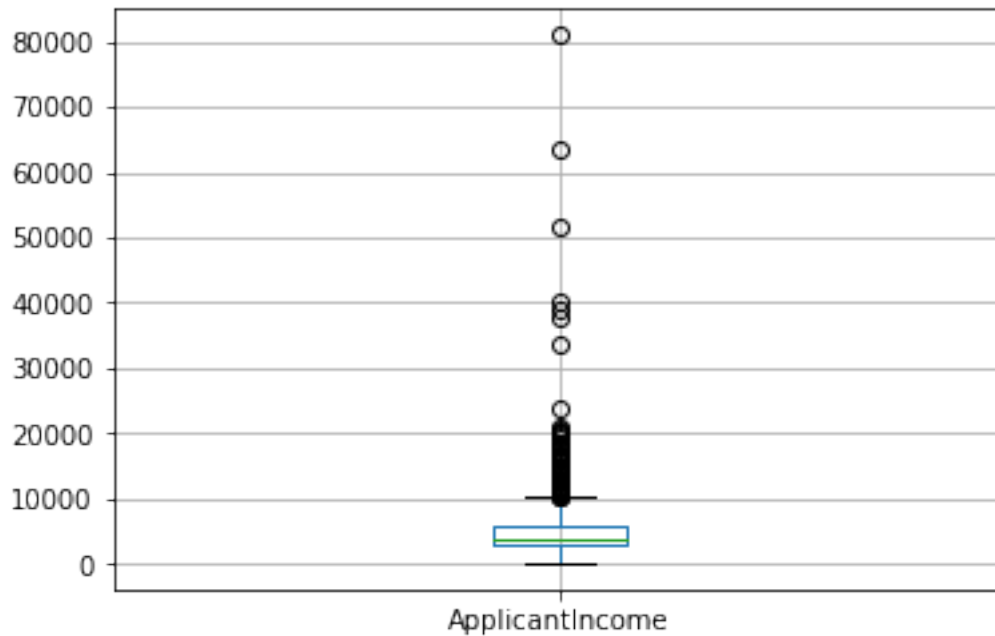
```
Out[53]: <matplotlib.axes._subplots.AxesSubplot at 0x7f93bc932780>
```



```
In [54]: # Box Plot for variable ApplicantIncome of training data set
```

```
df.boxplot(column='ApplicantIncome')
```

```
Out[54]: <matplotlib.axes._subplots.AxesSubplot at 0x7f93bc85e278>
```

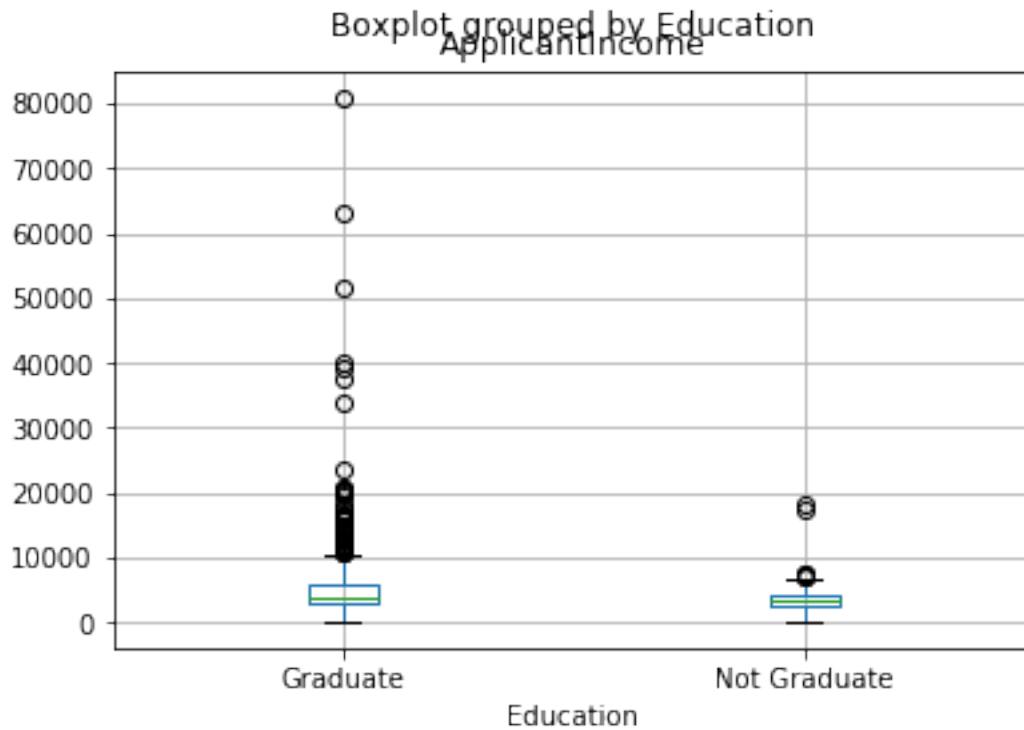


3. The above Box Plot confirms the presence of a lot of outliers/extreme values. This can be attributed to the income disparity in the society.

In [55]: *# Box Plot for variable ApplicantIncome by variable Education of training data set*

```
df.boxplot(column='ApplicantIncome', by = 'Education')
```

Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x7f93bc82e588>

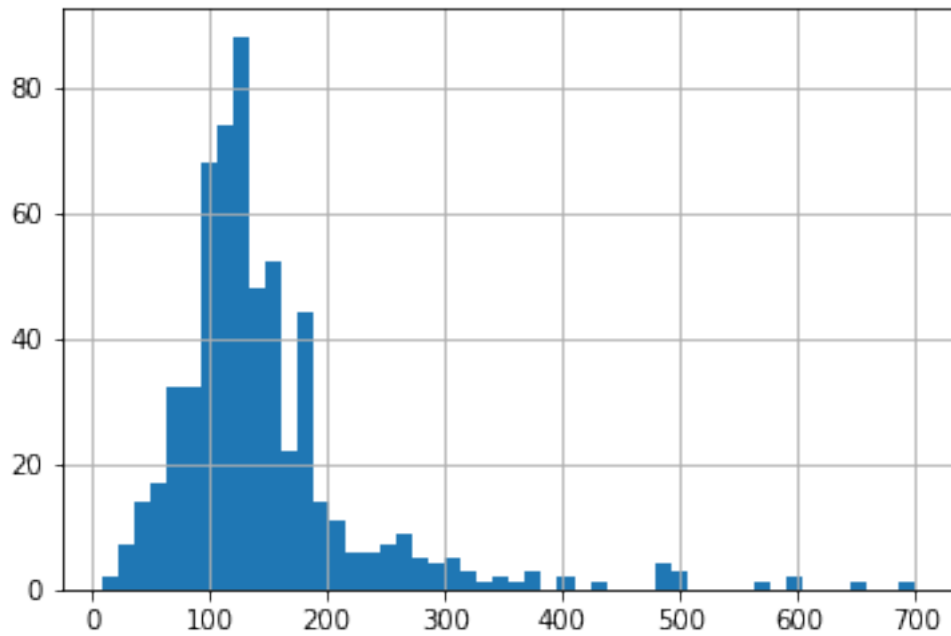


4. We can see that there is no substantial difference between the mean income of graduate and non-graduates. But there are a higher number of graduates with very high incomes, which are appearing to be the outliers

In [56]: # Histogram of variable LoanAmount

```
df['LoanAmount'].hist(bins=50)
```


Out[56]: <matplotlib.axes._subplots.AxesSubplot at 0x7f93bc73e2e8>



In [57]: # Box Plot for variable *LoanAmount* of training data set

```
df.boxplot(column='LoanAmount')
```

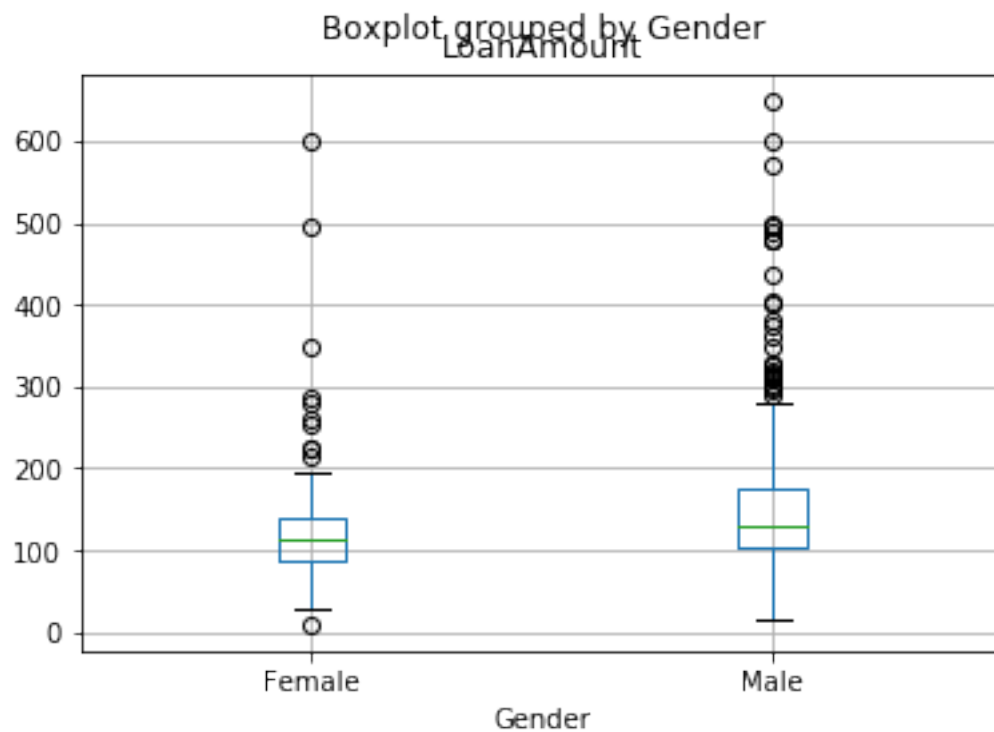
Out[57]: <matplotlib.axes._subplots.AxesSubplot at 0x7f93bc728be0>



```
In [58]: # Box Plot for variable LoanAmount by variable Gender of training data set
```

```
df.boxplot(column='LoanAmount', by = 'Gender')
```

```
Out[58]: <matplotlib.axes._subplots.AxesSubplot at 0x7f93bc79acc0>
```



5. LoanAmount has missing as well as extreme values, while ApplicantIncome has a few extreme values.

3 Understanding Distribution of Categorical Variables

```
In [15]: # Loan approval rates in absolute numbers
loan_approval = df['Loan_Status'].value_counts()['Y']
print(loan_approval)
```

422

- 422 number of loans were approved.


```
In [37]: # Credit History and Loan Status
pd.crosstab(df ['Credit_History'], df ['Loan_Status'], margins=True)
```

```
Out[37]: Loan_Status      N      Y  All
Credit_History
0.0           82      7   89
1.0           97    378  475
All          179    385  564
```

```
In [204]: #Function to output percentage row wise in a cross table
def percentageConvert(ser):
    return ser/float(ser[-1])

# Loan approval rate for customers having Credit_History (1)
df=pd.crosstab(df ["Credit_History"], df ["Loan_Status"], margins=True).apply(percentageConvert, axis=1)
loan_approval_with_Credit_1 = df['Y'][1]
print(loan_approval_with_Credit_1*100)
```

```
79.04761904761905
```

- 79.58 % of the applicants whose loans were approved have Credit_History equals to 1.

```
In [39]: df['Y']
```

```
Out[39]: Credit_History
0.0      0.078652
1.0      0.795789
All      0.682624
Name: Y, dtype: float64
```

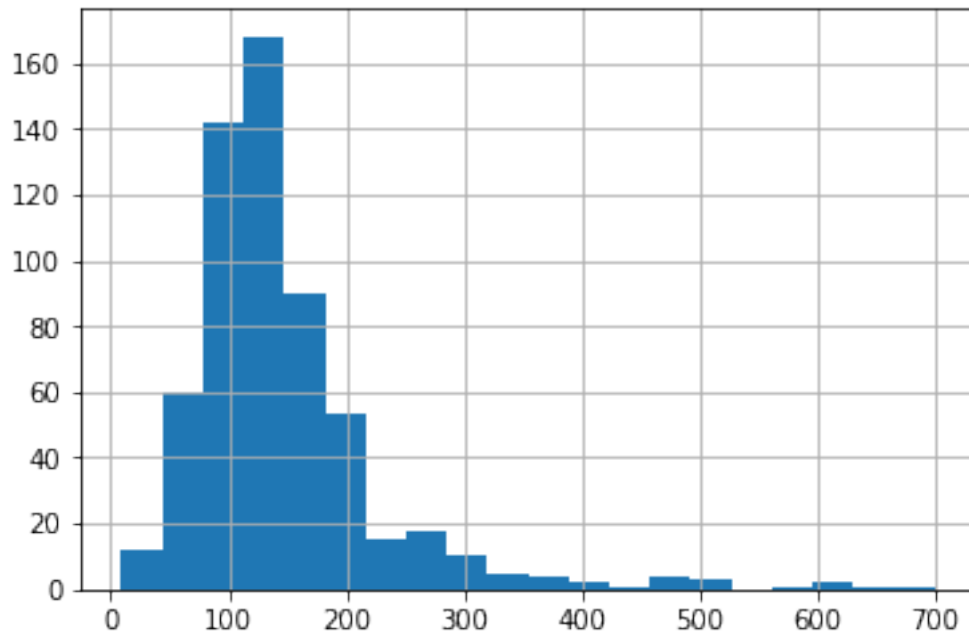
```
In [591]: # Replace missing value of Self_Employed with more frequent category
df['Self_Employed'].fillna('No', inplace=True)
```

4 Outliers of LoanAmount and Applicant Income

```
In [588]: # Add both ApplicantIncome and CoapplicantIncome to TotalIncome
df['TotalIncome'] = df['ApplicantIncome'] + df['CoapplicantIncome']

# Looking at the distribution of TotalIncome
df['LoanAmount'].hist(bins=20)
```

```
Out[588]: <matplotlib.axes._subplots.AxesSubplot at 0x7f6fad7ff98>
```

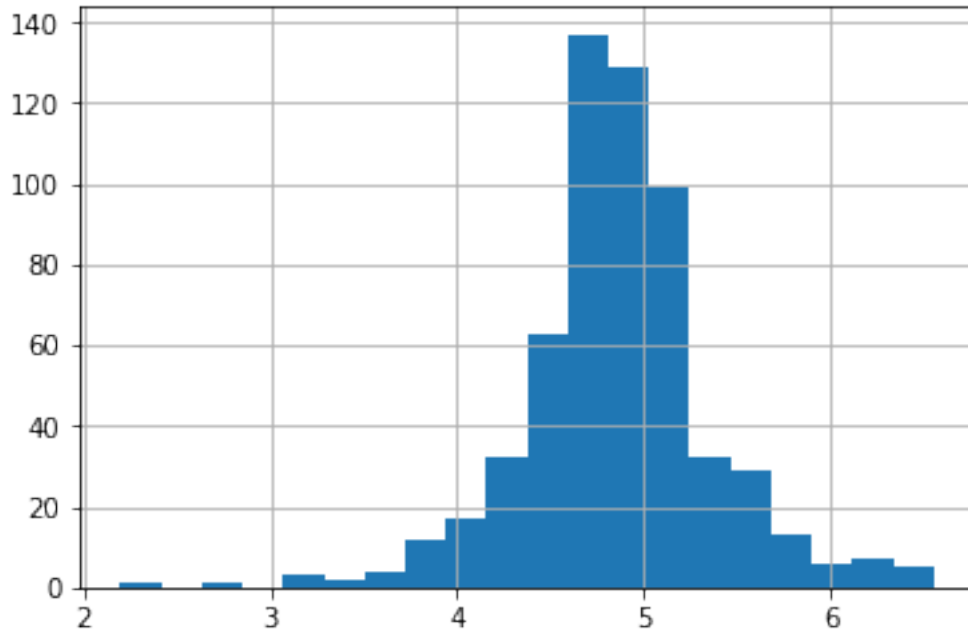


- The extreme values are practically possible, i.e. some people might apply for high value loans due to specific needs. So instead of treating them as outliers, let's try a log transformation to nullify their effect:

```
In [112]: # Perform log transformation of TotalIncome to make it closer to normal
          df['LoanAmount_log'] = np.log(df['LoanAmount'])

          # Looking at the distribution of TotalIncome_log
          df['LoanAmount_log'].hist(bins=20)
```

```
Out[112]: <matplotlib.axes._subplots.AxesSubplot at 0x7f93bbecec50>
```



5 Data Preparation for Model Building

- sklearn requires all inputs to be numeric, we should convert all our categorical variables into numeric by encoding the categories. Before that we will fill all the missing values in the dataset.

```
In [62]: # Impute missing values for Gender
df['Gender'].fillna(df['Gender'].mode()[0], inplace=True)

# Impute missing values for Married
df['Married'].fillna(df['Married'].mode()[0], inplace=True)

# Impute missing values for Dependents
df['Dependents'].fillna(df['Dependents'].mode()[0], inplace=True)

# Impute missing values for Credit_History
df['Credit_History'].fillna(df['Credit_History'].mode()[0], inplace=True)

# Convert all non-numeric values to number
cat=['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'Credit_History', 'Prop

for var in cat:
    le = preprocessing.LabelEncoder()
    df[var]=le.fit_transform(df[var].astype('str'))
df.dtypes
```

```
Out[62]: Loan_ID      object
        Gender        int64
        Married       int64
        Dependents    int64
        Education     int64
        Self_Employed int64
        ApplicantIncome int64
        CoapplicantIncome float64
        LoanAmount    float64
        Loan_Amount_Term float64
        Credit_History int64
        Property_Area int64
        Loan_Status   object
dtype: object
```

6 Generic Classification Function

```
In [208]: #Import models from scikit learn module:
          from sklearn import metrics
          from sklearn.cross_validation import KFold

          #Generic function for making a classification model and accessing performance:

          def classification_model(model, data, predictors, outcome):
              #Fit the model:
              model.fit(data[predictors],data[outcome])

              #Make predictions on training set:
              predictions = model.predict(data[predictors])

              #Print accuracy
              accuracy = metrics.accuracy_score(predictions,data[outcome])
              print ("Accuracy : %s" % "{0:.3%}".format(accuracy))

              #Perform k-fold cross-validation with 5 folds
              kf = KFold(data.shape[0], n_folds=5)
              error = []
              for train, test in kf:
                  # Filter training data
                  train_predictors = (data[predictors].iloc[train,:])

                  # The target we're using to train the algorithm.
                  train_target = data[outcome].iloc[train]

                  # Training the algorithm using the predictors and target.
                  model.fit(train_predictors, train_target)
```

```

        #Record error from each cross-validation run
        error.append(model.score(data[predictors].iloc[test,:], data[outcome].iloc[test,:]))

    print ("Cross-Validation Score : %s" % "{0:.3%}".format(np.mean(error)))

    #Fit the model again so that it can be referred outside the function:
    model.fit(data[predictors],data[outcome])

```

7 Model Building

In [186]: #Combining both train and test dataset

```

#Create a flag for Train and Test Data set
df['Type']='Train'
test['Type']='Test'
fullData = pd.concat([df,test],axis=0, sort=True)

#Look at the available missing values in the dataset
fullData.isnull().sum()

```

```

Out[186]: ApplicantIncome      0
CoapplicantIncome      0
Credit_History      29
Dependents      10
Education      0
Gender      11
LoanAmount      27
LoanAmount_log      389
Loan_Amount_Term      20
Loan_ID      0
Loan_Status      367
Married      0
Property_Area      0
Self_Employed      23
Type      0
dtype: int64

```

In [187]: #Identify categorical and continuous variables

```

ID_col = ['Loan_ID']
target_col = ["Loan_Status"]
cat_cols = ['Credit_History', 'Dependents', 'Gender', 'Married', 'Education', 'Property_Area']

```

In [200]: #Imputing Missing values with mean for continuous variable

```

fullData['LoanAmount'].fillna(fullData['LoanAmount'].mean(), inplace=True)
fullData['LoanAmount_log'].fillna(fullData['LoanAmount_log'].mean(), inplace=True)
fullData['Loan_Amount_Term'].fillna(fullData['Loan_Amount_Term'].mean(), inplace=True)
fullData['ApplicantIncome'].fillna(fullData['ApplicantIncome'].mean(), inplace=True)
fullData['CoapplicantIncome'].fillna(fullData['CoapplicantIncome'].mean(), inplace=True)

```

```

#Imputing Missing values with mode for categorical variables
fullData['Gender'].fillna(fullData['Gender'].mode()[0], inplace=True)
fullData['Married'].fillna(fullData['Married'].mode()[0], inplace=True)
fullData['Dependents'].fillna(fullData['Dependents'].mode()[0], inplace=True)
fullData['Loan_Amount_Term'].fillna(fullData['Loan_Amount_Term'].mode()[0], inplace=True)
fullData['Credit_History'].fillna(fullData['Credit_History'].mode()[0], inplace=True)

```

In [202]: #Create a new column as Total Income

```

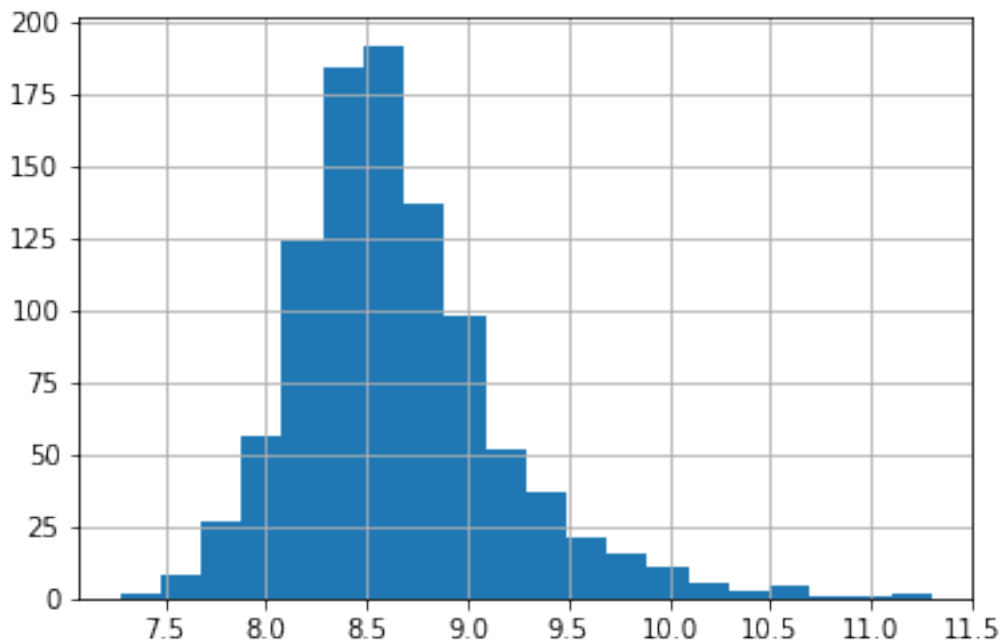
fullData['TotalIncome']=fullData['ApplicantIncome'] + fullData['CoapplicantIncome']

fullData['TotalIncome_log'] = np.log(fullData['TotalIncome'])

#Histogram for Total Income
fullData['TotalIncome_log'].hist(bins=20)

```

Out[202]: <matplotlib.axes._subplots.AxesSubplot at 0x7f93bbd93a20>



In [197]: #create label encoders for categorical features

```

for var in cat_cols:
    number = LabelEncoder()
    fullData[var] = number.fit_transform(fullData[var].astype('str'))

train_modified=fullData[fullData['Type']=='Train']
test_modified=fullData[fullData['Type']=='Test']
train_modified["Loan_Status"] = number.fit_transform(train_modified["Loan_Status"].ast

```

```
/home/parths007/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

7.1 Logistic Regression Model

1. The chances of getting a loan will be higher for:

- Applicants having a credit history (we observed this in exploration.)
- Applicants with higher applicant and co-applicant incomes
- Applicants with higher education level
- Properties in urban areas with high growth perspectives

So let's make our model with 'Credit_History', 'Education' & 'Gender'

```
In [198]: from sklearn.linear_model import LogisticRegression
```

```
predictors_Logistic=['Credit_History', 'Education', 'Gender']
```

```
x_train = train_modified[list(predictors_Logistic)].values
```

```
y_train = train_modified["Loan_Status"].values
```

```
x_test=test_modified[list(predictors_Logistic)].values
```

```
In [203]: # Create logistic regression object
```

```
model = LogisticRegression()
```

```
# Train the model using the training sets
```

```
model.fit(x_train, y_train)
```

```
#Predict Output
```

```
predicted= model.predict(x_test)
```

```
#Reverse encoding for predicted outcome
```

```
predicted = number.inverse_transform(predicted)
```

```
#Store it to test dataset
```

```
test_modified['Loan_Status']=predicted
```

```
outcome_var = 'Loan_Status'
```

```
classification_model(model, df, predictors_Logistic, outcome_var)
```

```
test_modified.to_csv("Logistic_Prediction.csv", columns=['Loan_ID', 'Loan_Status'])
```

Accuracy : 80.945%
Cross-Validation Score : 80.946%

```
/home/parths007/anaconda3/lib/python3.7/site-packages/sklearn/preprocessing/label.py:151: DeprecationWarning:
  if diff:
/home/parths007/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:14: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>